

AURA Technical Blueprint v1

Этот документ отвечает на вопрос: как именно собирать AURA MVP технически, сколько это примерно стоит при росте пользователей и какие решения нельзя откладывать до разработки. Он не заменяет финальное ТЗ, но задает архитектурную рамку, стек, unit economics и границы MVP.

Главное решение: MVP должен быть mobile-first, image-first и analytics-first. Видеоаватар, marketplace, community и тяжелый AI companion не входят в MVP, потому что главный риск сейчас - не наличие рынка, а прохождение петли Episode -> Action -> Reset -> Avatar -> Return Tomorrow.

1. Architecture Decision

Слой	Технология	Стоимость	Почему выбрана	Риск
Frontend	React Native + Expo	\$0 platform cost; developer time	Быстрее собрать iOS/Android MVP, легче подключить RevenueCat, push, analytics и OTA updates.	Сложные native animations и avatar/video editing могут потребовать native modules.
Backend	NestJS on Fly.io/Render/Railway or managed container	\$25-300/мо MVP, выше при росте	Четкая модульная архитектура, TypeScript, удобно для API, jobs, billing webhooks и admin.	Нужно дисциплинированно не строить enterprise backend раньше времени.
Database	Postgres via Supabase Pro	От \$25/мо; Pro включает 100k MAU и базовые лимиты	Postgres лучше подходит для связанных сущностей: user, season, episode, action, reflection, avatar state.	При тяжелых assets и логах нельзя складывать все в Postgres.
Storage	S3-compatible storage / Supabase Storage for MVP	Supabase Pro включает 100GB storage, далее usage-based	Нужно хранить avatar/canvas assets, recaps, exports.	Изображения быстро раздувают storage/egress, нужна компрессия и lifecycle policy.
AI Brain	OpenAI GPT-4.1 mini first, Claude/Gemini as fallback later	GPT-4.1 mini: \$0.40 input / \$1.60 output за 1M tokens	Хороший баланс цены, скорости, structured output и tool calling для MVP.	Качество "личности" может требовать prompt QA или более дорогой модели на отдельных шагах.
Image	Together FLUX.2 pro/dev or Replicate FLUX	Около \$0.015-\$0.04/image depending provider/model	Для MVP достаточно image-first Life Canvas; видео не нужно в core loop.	Если делать изображение каждый день для всех, себестоимость быстро становится главной статьей.
Video avatar	Not MVP; later Replicate/Tavus/HeyGen tests	Replicate video examples: \$0.09/sec 480p, \$0.25/sec 720p for Wan i2v	Видео может стать premium/token moment, но разрушает экономику free MVP.	Deepfake/privacy, latency, moderation, cost.
Payments	RevenueCat + App Store/Google Play IAP	Free up to \$2.5k MTR, then 1% tracked revenue; store fee separately	Быстрый subscription stack, restore purchases, webhooks, entitlement source of truth.	Нужно считать store fee и RevenueCat fee отдельно.
Analytics	PostHog / Amplitude / Firebase Analytics	\$0-\$300/мо early depending tool and events	MVP без analytics бессмысленен: нужно видеть loop completion, D1, D7, paywall.	Слишком много событий без product questions.
Admin	Retool/Supabase Studio/internal Next.js later	\$0-\$50/мо early + dev time	Нужно менять prompts, season templates, смотреть safety flags без релиза приложения.	Не строить большой backoffice до первых пользователей.

2. System Architecture

Блок	Что делает	Интеграции	Данные
Mobile App	Onboarding, profile, season, episode, action, reset, reflection, Life Canvas, paywall, reminders.	RevenueCat SDK, analytics SDK, push, backend API.	Сохраняет минимум локального cache, не хранит secrets.
API Backend	Auth session, profile, seasons, episodes, actions, avatar jobs, billing webhooks, admin.	Supabase/Postgres, OpenAI, image provider, RevenueCat, analytics.	Единый слой бизнес-правил и state transitions.
Postgres	Хранит связанные сущности продукта.	Supabase APIs, backend service role.	UserProfile, Season, Episode, Action, Reflection, AvatarState, Subscription, Event.
AI Layer	Генерирует episode/action/reset recap, structured JSON, safety checked outputs.	OpenAI primary, optional Claude/Gemini fallback.	PromptVersion, request/response metadata, cost logs.
Image Layer	Генерирует image-first Life Canvas shifts.	Together/Replicate/other provider.	Prompt, provider, asset id, latency, cost.
Storage	Хранит изображения, exports, preview cards.	S3/Supabase Storage/CDN.	Compressed assets, thumbnails, lifecycle policy.
Billing	Entitlements, trial, subscription status, webhooks.	RevenueCat, App Store, Google Play.	Plan, status, renewal, cancellation reason.
Analytics	Funnel, retention, paywall, cost-per-loop, avatar causality.	PostHog/Amplitude/Firebase, warehouse later.	Events with user_state and experiment_id.
Admin	Prompt publish, safety review, season templates, user issue review.	Backend admin endpoints.	Prompt versions, moderation queue, feature flags.

3. Recommended Stack For MVP

Рекомендация: React Native + Expo, NestJS, Supabase Postgres, S3/Supabase Storage, OpenAI GPT-4.1 mini, Together/Replicate FLUX, RevenueCat, PostHog/Firebase Analytics, Retool or simple internal admin. Это не “идеальный навсегда” стек, а стек, который быстрее всего доводит продукт до интервью, прототипа и первых 30-100 пользователей.

Решение	Почему	Альтернатива
React Native вместо Flutter	Больше готовых integrations для paywall/analytics/AI tooling и быстрее нанимать JS/TS команду.	Flutter тоже валиден, если команда уже сильнее во Flutter.
NestJS вместо no-code backend	Нужны state transitions, async jobs, billing webhooks, admin endpoints и cost logs.	Supabase Edge Functions можно использовать для быстрого старта, но бизнес-логика быстро усложнится.
Postgres вместо Firestore	У AURA много связанных сущностей и отчетности; SQL проще для аналитики и unit economics.	Firebase хорош для realtime/simple mobile apps, но relational product graph здесь важнее.

Решение	Почему	Альтернатива
Image-first вместо video-first	Видео дорого, медленно и рискованно; центральную гипотезу avatar causality можно проверить изображением.	Видео тестировать как premium token после proof of loop.
One primary LLM	На MVP лучше контролировать prompt quality и стоимость, чем строить router.	OpenRouter/LLM router после 1000+ активных пользователей и накопления evals.
RevenueCat с первого paid test	Restore, entitlements, webhooks и paywall experiments не стоит писать руками.	Чистый StoreKit/Google Billing только если команда хочет контролировать billing сама.

4. Component Responsibilities

Компонент	Ответственность	Must	Do not
Mobile Shell	Navigation, auth state, onboarding, paywall routing, offline event queue.	Fast cold start, no blank loading before Welcome.	Do not embed provider secrets or prompt logic.
Onboarding Module	Privacy, consent, birth date, current state, season choice.	Can resume if app closes mid-flow.	Do not ask 20 questions before first episode.
Episode Module	Render episode, read progress, dislike/report, action entry.	Supports generated and fallback episodes.	Do not expose raw AI output without product formatting.
Action Module	Action choices, difficulty, reminders, completion state.	Always offer a softer version.	Do not punish skipped or partial action.
Reset Module	Short ritual, timer, skip, complete, event logging.	30-60 sec default.	Do not become a full meditation product in MVP.
Reflection Module	Emotion after, one-line note, save to memory.	Can save without text.	Do not require journaling friction.
Life Canvas Module	Pending state, generated image, explanation, save/share.	Show why canvas changed.	Do not show random beauty without causality.
Memory Module	Season history, recap, locked paid sections.	Make 7-day trajectory visible.	Do not overbuild timeline/social feed.
Billing Module	Paywall config, trial, purchase, restore, entitlement.	Single entitlement source of truth.	Do not hardcode prices in app.
Admin Module	Prompt versions, season templates, safety queue, manual overrides.	Prompt rollback.	Do not build a complex CRM.

5. Database Schema Draft

Table	Core columns	Indexes	Notes
users	id, email_hash, auth_provider, locale, timezone, created_at, deleted_at	id, auth_provider	PII минимум; не хранить лишние персональные данные.
user_profiles	user_id, display_name, birth_date, current_goal, mood, privacy_flags, updated_at	user_id	Birth date is sensitive; access only through backend.

Table	Core columns	Indexes	Notes
consents	id, user_id, policy_version, accepted_at, revoked_at	user_id, policy_version	Нужно для доверия и compliance.
season_templates	id, theme, title, description, active, sort_order	active	Управляется из admin.
seasons	id, user_id, template_id, status, day_index, started_at, completed_at	user_id, status	Один active season в MVP.
prompt_versions	id, type, version, template, safety_rules, active, created_at	type, active	Все AI outputs должны знать prompt version.
episodes	id, season_id, day_index, title, insight, conflict, resource, risk, prompt_version_id, safety_status	season_id, day_index	Structured JSON, не raw blob only.
actions	id, episode_id, difficulty, text, estimated_minutes, status, selected_at, completed_at	episode_id, status	Action is observable behavior.
reset_sessions	id, action_id, reset_type, duration_sec, status, started_at, completed_at	action_id	Skip allowed.
reflections	id, action_id, emotion_after, note, created_at	action_id	One-line reflection.
avatar_states	id, user_id, season_id, episode_id, action_id, visual_traits_json, explanation, asset_id, created_at	user_id, season_id	Core causal visual state.
assets	id, user_id, provider, type, url, thumbnail_url, status, estimated_cost, latency_ms	user_id, provider, status	Cost and provider logging mandatory.
subscriptions	user_id, store, revenuecat_id, entitlement, status, trial_start, renewal_at	user_id, status	RevenueCat webhook source.
notifications	id, user_id, type, scheduled_at, sent_at, opened_at, status	user_id, scheduled_at	D1/D7 return loop.
events	id, user_id, event_name, user_state, properties_json, client_ts, server_ts	event_name, server_ts, user_id	Could later move to warehouse.
generation_logs	id, user_id, provider, model, input_units, output_units, estimated_cost, latency_ms, status	provider, model, created_at	Unit economics depends on this.

6. API Groups

Group	Endpoints	Owner	QA
Auth/Profile	POST /auth/session, PATCH /profile, GET /profile, DELETE /profile	Backend/mobile	Can create, resume, export/delete data.
Consent	GET /consent/current, POST /consent/accept, POST /consent/revoke	Backend	Policy version stored.
Season	GET /seasons/templates, POST /seasons, GET /seasons/current, PATCH /seasons/:id	Backend/product	Only one active MVP season.

Group	Endpoints	Owner	QA
Episode	POST /episodes/generate, GET /episodes/today, POST /episodes/:id/report	AI/backend	Fallback and safety status.
Action	POST /actions/select, PATCH /actions/:id, POST /actions/:id/complete	Backend/mobile	Partial completion valid.
Reset	POST /reset/start, POST /reset/:id/complete	Backend/mobile	Skip does not break flow.
Reflection	POST /reflections, PATCH /reflections/:id	Backend/mobile	Text optional.
Avatar	POST /avatar/generate, GET /avatar/:job_id, GET /avatar/current	AI/image/backend	Async job, pending/failure/retry.
Memory	GET /memory, GET /recap/weekly, POST /recap/generate	Backend/AI	Free locks and paid unlocks.
Billing	GET /paywall, POST /billing/webhook, GET /entitlements	Backend/mobile	Sandbox purchase and restore.
Events	POST /events, POST /events/batch	Data/mobile	Offline batch and dedupe.
Admin	GET/POST /admin/prompts, /admin/templates, /admin/moderation	Backend/internal	Role protected, audit log.

7. Provider Comparison

Category	Provider	Pros	Cons	Verdict
LLM primary	OpenAI GPT-4.1 mini	Low cost, strong structured output, good speed.	May need prompt tuning for warmth.	Use for MVP default.
LLM quality fallback	Claude Sonnet	Strong writing/reasoning quality.	Much more expensive per token.	Use only for evals or premium generation later.
LLM low-cost fallback	Gemini/DeepSeek via router	Potentially cheaper/fast.	Quality and safety variance.	Not day-one unless team already has eval harness.
Image primary	Together FLUX.2 dev/pro	Clear per-image pricing and API.	Style consistency needs prompts/seeds.	Good MVP candidate.
Image alternative	Replicate FLUX	Many models and fast experimentation.	Costs vary by model; latency varies.	Good for experiments and fallback.
Video later	Replicate/Tavus/HeyGen	Can create premium wow moments.	Expensive, privacy/deepfake risk, longer latency.	Not free MVP.
Billing	RevenueCat	Fast subscriptions, entitlements, webhooks.	Adds fee after revenue threshold.	Use from first paid test.
Analytics	PostHog/Firebase/Amplitude	Funnels, cohorts, events.	Pricing can rise with events.	Use one; do not duplicate event stacks.

8. Unit Economics Assumptions

Ниже модель не претендует на бухгалтерскую точность. Ее задача - понять порядок величин и главный риск. Допущение нормального MVP: 1 MAU в среднем имеет 8 активных дней в месяц, получает 4 image-based avatar shifts в месяц, daily AI использует GPT-4.1 mini, видеоаватар не входит в MVP. Store fee, VAT/tax, paid acquisition и зарплаты команды не включены в product COGS и считаются отдельно.

Статья	Модельное допущение	Источник / основание
AI text	8 active days/user/month; ~3,200 input tokens and ~1,000 output tokens per active day	OpenAI GPT-4.1 mini public pricing: \$0.40 input and \$1.60 output per 1M tokens.
Image	4 images/user/month at \$0.03 each	Together FLUX.2 pro around \$0.03/image; Replicate FLUX ranges around \$0.003-\$0.04/image depending model.
Storage	\$0.006/user/month placeholder	Supabase Pro includes 100GB file storage, then overage.
Analytics	\$0.01/user/month placeholder	Tool-dependent; MVP can start low/free but events must be planned.
Push	\$0.002/user/month placeholder	Firebase/APNs costs usually not the first bottleneck for MVP.
Support	\$0.025/user/month placeholder	Internal support/ops placeholder for early cohorts.
Infra base	\$65 at 100-1,000 MAU; \$220 at 10k; \$1,300 at 100k	Supabase/hosting/monitoring/admin estimates, not exact invoices.
Optional video	15% of MAU get one 5-sec video at \$0.09/sec in stress scenario	Replicate video example \$0.09/sec 480p.

9. Unit Economics By Scale

MAU	AI text	Images	Storage	Analytics	Push	Support	Infra/Admin	Total/mo	Cost/MAU	If optional video
100	\$2	\$12	\$1	\$1	\$0	\$3	\$65	\$84	\$0.84	\$90
1000	\$23	\$120	\$6	\$10	\$2	\$25	\$65	\$251	\$0.25	\$319
10000	\$230	\$1,200	\$60	\$100	\$20	\$250	\$220	\$2,080	\$0.21	\$2,755
100000	\$2,304	\$12,000	\$600	\$1,000	\$200	\$2,500	\$1,300	\$19,904	\$0.20	\$26,654

Главный вывод: в image-first MVP себестоимость не выглядит убийственной, но изображения уже становятся основной переменной статьей. Видео нельзя давать всем бесплатно: даже один 5-секундный video moment для 15% MAU заметно увеличивает расходы. Поэтому video avatar должен быть либо premium/token, либо validation-only ручным экспериментом.

10. Sensitivity Analysis

Variable	Low	Base	High	Impact
Images per MAU	2/mo	4/mo	12/mo	Главный variable cost. Daily images for free users can triple cost.
Image model	\$0.003/image	\$0.03/image	\$0.04+/image	Provider/model choice can change image COGS by 10x.
Active days	4/mo	8/mo	20/mo	Retention increases value but also AI usage.
Paid conversion	1%	5%	10%	At low conversion free COGS must be aggressively limited.
Video usage	0	premium only	free daily	Free daily video likely breaks MVP margin.

Variable	Low	Base	High	Impact
Store fee	15%	30%	30% + taxes	Subscription net revenue must be modeled after platform fees.
Support load	self-serve	light	high-touch wellbeing	If users treat AURA like coaching/therapy, support cost and safety burden rise.

11. Revenue And Margin Scenarios

MAU	Paid conversion	Paid users	Net revenue/mo after 30% store fee at \$7.99	Product COGS/mo	Вывод
1,000	1%	10	\$56	\$251	Убыточно без ограничения free costs. 1% paid conversion недостаточен.
1,000	5%	50	\$280	\$251	Почти бьется product COGS, но не покрывает маркетинг и команду.
10,000	3%	300	\$1,678	\$2,080	Нужно либо лучше конвертировать, либо снижать image frequency/free cost.
10,000	7%	700	\$3,914	\$2,080	Product gross margin начинает выглядеть рабочей, paid acquisition еще не доказан.
100,000	5%	5,000	\$27,965	\$19,904	Становится бизнесом только при сильном retention и контроле image/video costs.

12. Cost Control Rules

Правило	Почему	Как внедрить
No free daily video	Видео может стоить больше подписочной маржи.	Video only paid token/premium validation.
Image budget per free user	Images dominate variable COGS.	Free: 1-2 images/week; paid: more styles and recaps.
Prompt caching and templates	Большая часть system prompt повторяется.	Cache where provider supports it; use structured templates.
Async generation	Latency and retries не должны блокировать UX.	Queue, pending state, retry, fallback.
Cost log per generation	Без cost-per-loop нельзя принимать pricing decisions.	provider, tokens, images, seconds, latency, estimated_cost.
Manual QA before scaling	AI quality важнее автомасштаба до 100 users.	Prompt review and admin override for early cohorts.

13. Security And Privacy Requirements

Requirement	Reason	Implementation
PII minimization	Дата рождения чувствительна.	Не собирать адрес, телефон, фото лица в MVP unless absolutely needed.
Data deletion	Trust and compliance.	DELETE /profile and asset deletion job.
Consent versioning	Нужно знать, с какой политикой согласился user.	consents table with policy_version.
Provider secrecy	API keys and prompts are sensitive.	All AI calls via backend only.
Asset privacy	Avatar/Life Canvas может быть личным.	Private buckets, signed URLs, thumbnails.
Safety moderation	Wellbeing/spirituality adjacent risk.	Prompt rules, output filter, report issue, admin review.
No deterministic claims	Не обещать судьбу/диагноз.	AI output contract forbids medical/absolute predictions.
Audit log	Admin changes prompts and templates.	admin_actions table.

14. Build Phases

Фаза	Что строим	Цель	Условие выхода
Phase 0	Clickable prototype + fake backend + manual content	20 interviews and comprehension test.	70% understand Episode -> Action -> Avatar causality.
Phase 1	Real auth/profile/season/episode/action/reflection/image/paywall sandbox	30-100 concierge users.	Completed loop and D1 return measurable.
Phase 2	Notifications, memory, weekly recap, admin prompts, cost dashboard	1000 users without chaos.	D7 and trial intent visible.
Phase 3	Provider optimization, experiments, paid acquisition readiness	Scale only if unit economics and retention pass.	CAC/LTV test can start.

15. Engineering Roadmap

Sprint	Deliverable	Output	Risk
Sprint 1	Repo setup, app shell, backend skeleton, DB schema, auth/session.	User can open app, create session, accept consent.	Too much time on design polish before flow works.
Sprint 2	Profile, season templates, episode prompt v1, structured output.	Day 1 episode generated and saved.	AI output generic; needs prompt QA.
Sprint 3	Action, reset, reflection, analytics events.	Completed loop without avatar.	Action friction too high.
Sprint 4	Image provider, avatar state, storage, pending/failure states.	Life Canvas shift with explanation.	Visual causality unclear.
Sprint 5	Paywall sandbox, RevenueCat webhook, entitlements.	Trial/subscription unlocks memory/premium flags.	Paywall timing hurts D1.

Sprint	Deliverable	Output	Risk
Sprint 6	Memory, weekly recap, notifications.	7-day season can complete.	Recap weak without enough user data.
Sprint 7	Admin prompts/templates/safety queue.	Team can operate first cohort.	Admin scope creep.
Sprint 8	QA, dashboards, cost logging, soft launch.	30-100 users measured end-to-end.	No single owner for event quality.

16. Technical Risk Register

Risk	Probability	Impact	Mitigation
AI output feels generic	High	High	PromptVersion, manual QA, profile context, user feedback flag.
Image not causally understood	High	Very high	Explanation copy, visual trait mapping, prototype interviews.
Video cost explodes	Medium	High	No free video in MVP, premium/token only.
Paywall breaks trust	Medium	High	Paywall only after completed loop.
Data/privacy concerns	Medium	High	Consent, deletion, no face upload in MVP.
Analytics unreliable	Medium	High	Event taxonomy, QA checklist, dashboard before launch.
Admin not ready	Medium	Medium	Minimal prompt/template controls before 100 users.
Backend overbuilt	Medium	Medium	Build only core loop modules before community/marketplace.

17. Implementation Backlog

ID	Task	Area	Output	Priority
T-01	Create mobile app shell	Frontend	Navigation, session state, empty Welcome.	P0
T-02	Implement consent flow	Frontend/Backend	Policy version accepted and stored.	P0
T-03	Profile form	Frontend	Birth date, mood, current goal saved.	P0
T-04	Season templates API	Backend	Active templates returned to app.	P0
T-05	Start season endpoint	Backend	One active season per user.	P0
T-06	PromptVersion registry	AI/Admin	Prompt templates versioned.	P0
T-07	Episode generation endpoint	AI/Backend	Structured episode JSON saved.	P0
T-08	Episode safety filter	AI/Safety	No diagnosis/prediction/harmful advice.	P0
T-09	Episode screen	Frontend	Title, insight, conflict, resource, risk.	P0
T-10	Action options generation	AI/Backend	Easy/normal/brave actions.	P0
T-11	Action selection UI	Frontend	Select and soften action.	P0
T-12	Reset timer	Frontend	30-60 sec reset with skip/complete.	P0
T-13	Reflection UI	Frontend	Emotion + optional note.	P0
T-14	Avatar state data model	Backend	Causal visual traits saved.	P0

ID	Task	Area	Output	Priority
T-15	Image provider integration	AI/Image	Generate and store Life Canvas asset.	P0
T-16	Avatar pending/failure states	Frontend	No blocked UI during generation.	P0
T-17	Avatar explanation copy	Product/Frontend	Changed because of action X.	P0
T-18	Event collector	Analytics	POST /events and batch retry.	P0
T-19	Core funnel dashboard	Analytics	Activation, loop, D1, paywall.	P0
T-20	Cost logging	Backend/Data	Provider, tokens/images, estimated cost.	P0
T-21	RevenueCat sandbox	Billing	Purchase, restore, entitlement.	P0
T-22	Paywall config API	Backend	Placement and copy variants.	P1
T-23	Memory archive	Frontend/Backend	Season timeline.	P1
T-24	Weekly recap	AI/Frontend	7-day recap generated.	P1
T-25	Push notification setup	Mobile	Tomorrow hook reminder.	P1
T-26	Comeback flow	Product/Frontend	Missed day recovery.	P1
T-27	Admin prompt publish	Admin	Publish/rollback prompt versions.	P1
T-28	Admin safety queue	Admin/Safety	Review reported outputs.	P1
T-29	Data export/delete	Privacy	User can remove data.	P1
T-30	Share card	Growth/Frontend	Optional card after avatar shift.	P2
T-31	Style presets	Image/Product	2-3 Life Canvas styles.	P2
T-32	Provider fallback	AI/Backend	Switch provider on failure.	P2
T-33	LLM eval set	AI/Data	Manual examples and scoring.	P2
T-34	A/B experiment flags	Backend/Product	Paywall and copy variants.	P2
T-35	Video token experiment	AI/Image	Paid/manual only.	Later
T-36	Community features	Product	Not MVP.	Later

18. Event Taxonomy

Event	When	Properties	Decision
app_opened	Every app open.	source, user_state, app_version	Baseline activity.
onboarding_started	Welcome start.	source, campaign	Top funnel.
consent_accepted	Privacy accepted.	policy_version	Trust barrier.
profile_started	Profile screen opened.	fields_visible	Onboarding friction.
profile_completed	Profile saved.	fields_count, skipped_birth_date	Personalization readiness.
season_template_viewed	Season select opened.	recommended_theme	Theme demand.
season_started	Season chosen.	theme, day_index	Season intent.

Event	When	Properties	Decision
episode_generate_started	Backend generation starts.	model, prompt_version	Latency/cost.
episode_generated	Episode saved.	latency_ms, estimated_cost, safety_status	AI reliability.
episode_read	Read threshold reached.	read_time, scroll_depth	Content engagement.
episode_reported	User reports issue.	reason, prompt_version	Safety/product quality.
action_selected	Action chosen.	difficulty, estimated_minutes	Action fit.
action_softened	User asks easier version.	original_difficulty	Difficulty tuning.
reset_started	Reset begins.	type, duration	Reset adoption.
reset_completed	Reset complete.	duration, skipped=false	Ritual value.
action_completed	Action marked done.	status, time_to_complete	Core behavior.
reflection_saved	Emotion/note saved.	emotion_after, note_length	Memory friction.
avatar_generate_started	Image job starts.	provider, style	Image pipeline.
avatar_generated	Asset ready.	latency_ms, estimated_cost, provider	Visual cost/quality.
avatar_causality_checked	User answers/interaction shows understanding.	understood=true/false	Main hypothesis.
tomorrow_hook_seen	Hook shown.	day_index, reminder_prompted	Return setup.
push_opt_in	Notifications enabled.	placement	Reminder viability.
paywall_viewed	Paywall shown.	placement, offer, price	Monetization timing.
trial_started	Trial starts.	plan, store	WTP.
subscription_started	Paid starts.	plan, net_price	Revenue.
season_completed	Day 7 complete.	actions_completed, recaps_saved	D7 success.
share_card_clicked	Share card opened.	screen, style	Virality.
data_delete_requested	User requests deletion.	reason optional	Trust/privacy.

19. QA Checklist Before First Cohort

Area	Check	Owner
Onboarding	User can complete profile in under 90 seconds.	Product/QA
Privacy	Consent version stored and visible.	Backend
Episode	10 test profiles produce safe, non-generic outputs.	AI/Product
Action	Every episode has easy/normal/brave action.	AI/Product

Area	Check	Owner
Reset	Reset can complete, skip and resume without broken state.	Mobile
Reflection	Emotion-only save works.	Mobile/Backend
Avatar	Pending, success, failure and retry states work.	Mobile/Backend
Causality	Avatar explanation names the action that caused change.	Product
Storage	Assets are private and have thumbnails.	Backend
Billing	Sandbox purchase, restore and webhook sync work.	Mobile/Backend
Analytics	Core funnel events visible in dashboard.	Data
Cost	Every AI/image call writes estimated cost.	Backend/Data
Safety	Report issue creates admin review item.	Backend/Admin
Notifications	Tomorrow reminder can be scheduled and opened.	Mobile
Memory	Day 1 and Day 2 references connect correctly.	Backend/AI
Deletion	Data deletion path tested on staging.	Backend

20. Source Notes

Источник	Что взяли в модель	URL
OpenAI	GPT-4.1 mini pricing used: \$0.40 input / \$1.60 output per 1M tokens.	https://developers.openai.com/api/docs/models/gpt-4.1-mini
Anthropic	Claude Sonnet 4.6 pricing: \$3 input / \$15 output per 1M tokens; Haiku 4.5: \$1 / \$5.	https://platform.claude.com/docs/en/about-claude/pricing
Together AI	Image pricing examples include FLUX.2 pro around \$0.03/image and FLUX.2 dev around \$0.0154/image.	https://www.together.ai/pricing
Replicate	FLUX and video examples used for stress testing; e.g. FLUX 1.1 pro \$0.04/image and Wan i2v 480p \$0.09/sec.	https://replicate.com/pricing
Supabase	Pro starts at \$25/mo, includes 100k MAU, 8GB disk, 250GB egress, 100GB file storage.	https://supabase.com/pricing

Источник	Что взяли в модель	URL
RevenueCat	Starts free up to \$2,500 MTR, then 1% tracked revenue after that threshold.	https://www.revenuecat.com/pricing/

21. Final Technical Decision

Строить AURA MVP нужно как легкий мобильный продукт с жестким контролем AI/image costs. Архитектурно важно не переусложнить запуск: React Native, NestJS, Postgres, RevenueCat, один основной LLM, image-first avatar, простая admin-панель и сильная аналитика. Финансово важнейшее ограничение: free пользователь не должен получать дорогой ежедневный video/avatar pipeline. До доказательства D1/D7 и willingness to pay нужно считать не “сколько стоит разработать мечту”, а “сколько стоит один completed loop и один вернувшийся завтра пользователь”.